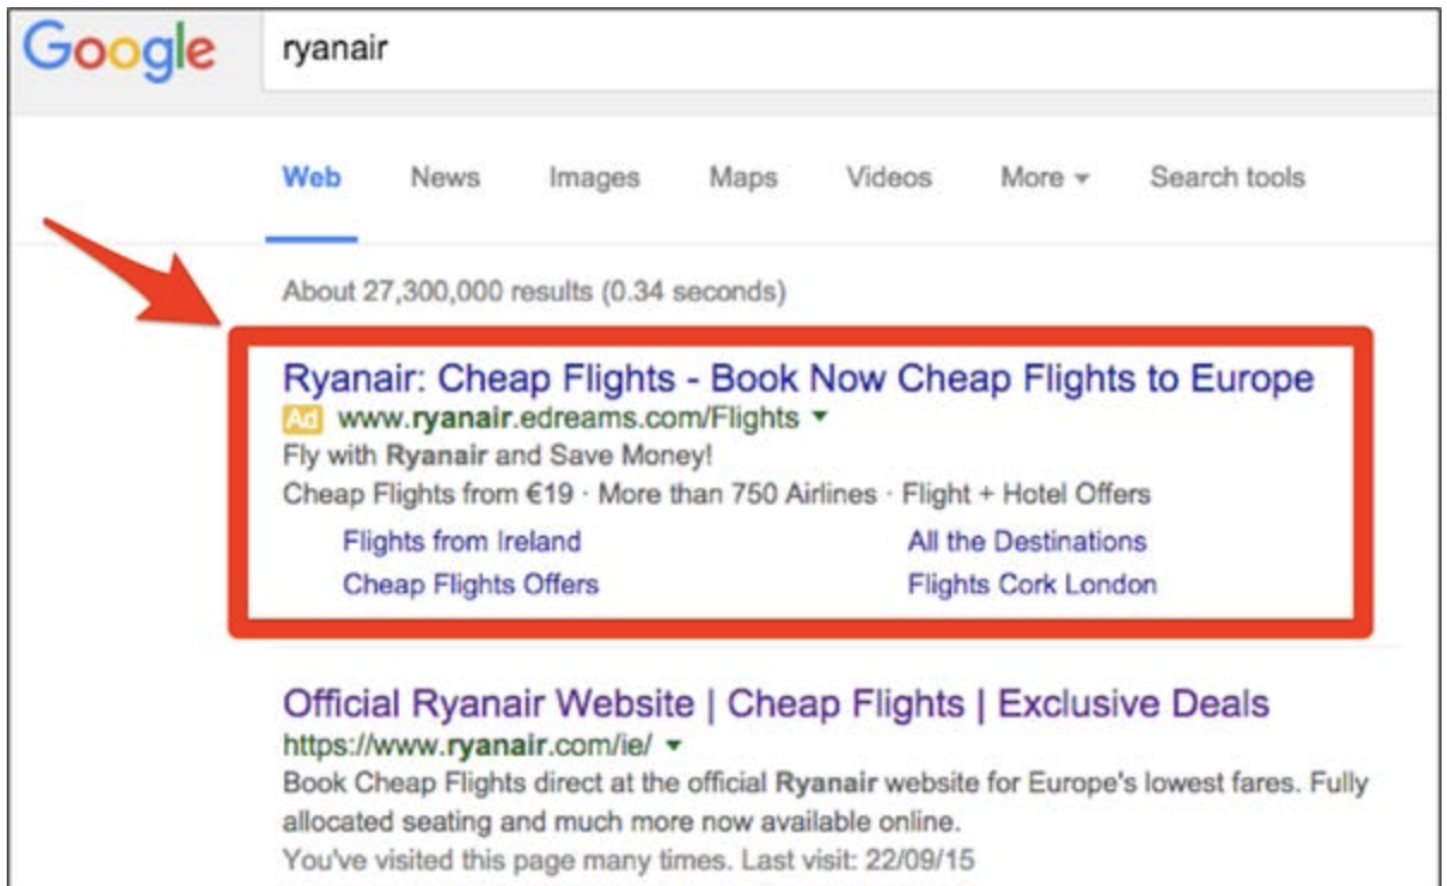


#programming #systemDesign


Advertisement Clicks Aggregator



Other Examples

- YouTube video clicks
- Instagram post clicks
- Website clicks
- App clicks

Simple System



- Simple enough design
- Will work for small applications

Will it scale?

Let's say you have **300M DAU**.

Each user has **10 meaningful clicks** every day

$300M * 10 = 3 \text{ Billion clicks / day}$


$3B / 24 / 60 / 60 = \sim 35K \text{ QPS}$

Peak QPS = $35K * 2 = 70K \text{ QPS}$

One click data is **0.1KB**

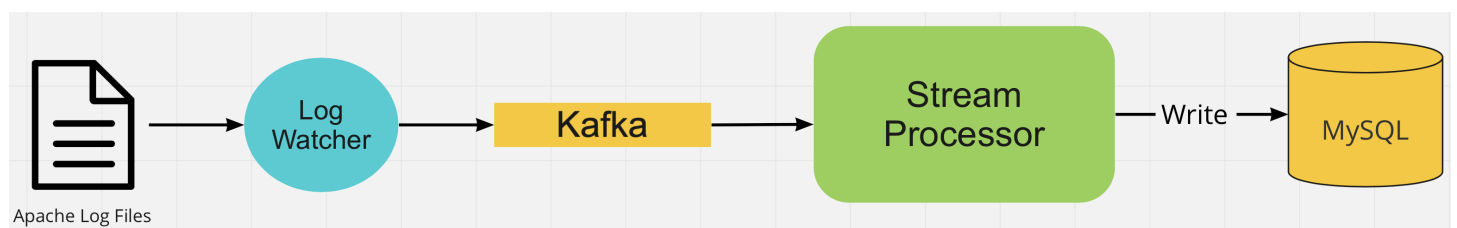
Each day click storage = $0.1KB * 3B = 3TB/day$

Issue 1: Server Overload



- High, Spiky QPS
- Server will need to be overprovisioned
- Synchronous API call more likely to run into errors

So we need to make the system more scalable, elastic & fault tolerant.



- All clicks come to log files
- Watcher to read new logs as they come along
- Writes a message to Kafka with all relevant details

- Lightweight stream processor consumes from Kafka and writes to MySQL

How can we scale Kafka?

- Add more partitions

How can we scale stream processor?

- Add more instances as you add more Kafka partitions
- Bump memory of instance

Issue 2: Storage

Our application requirements:

- High, spiky QPS
- Write intensive
- New data: 3TB/day
- Do we care about storing individual clicks?
 - We care about aggregations more
 - But it's important to keep individual clicks for historical purposes
- Read pattern?
 - Read aggregates such as:
 - Total clicks in last 20 minutes
 - Total clicks in last 30 days
 - Total clicks per advertisements in the last 90 day

Issue 1

High write throughput might introduce too much latency in the system.

Issues with replication lag.

We need a write intensive database.

Issue 2

3TB/day means lots of storage consumed per day.

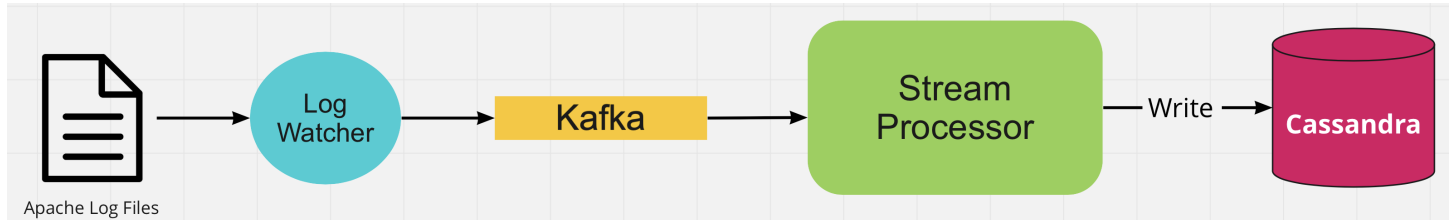
Storing so much data can be expensive.

Issue 3

We mainly care about aggregations, not individual clicks.
Relational DBs won't do well with aggregation queries.

Solution

Use Cassandra.



Cassandra features:

- Handles high write throughput well
- Horizontally scales, so can store huge volume of data easily
- Good for time-series data

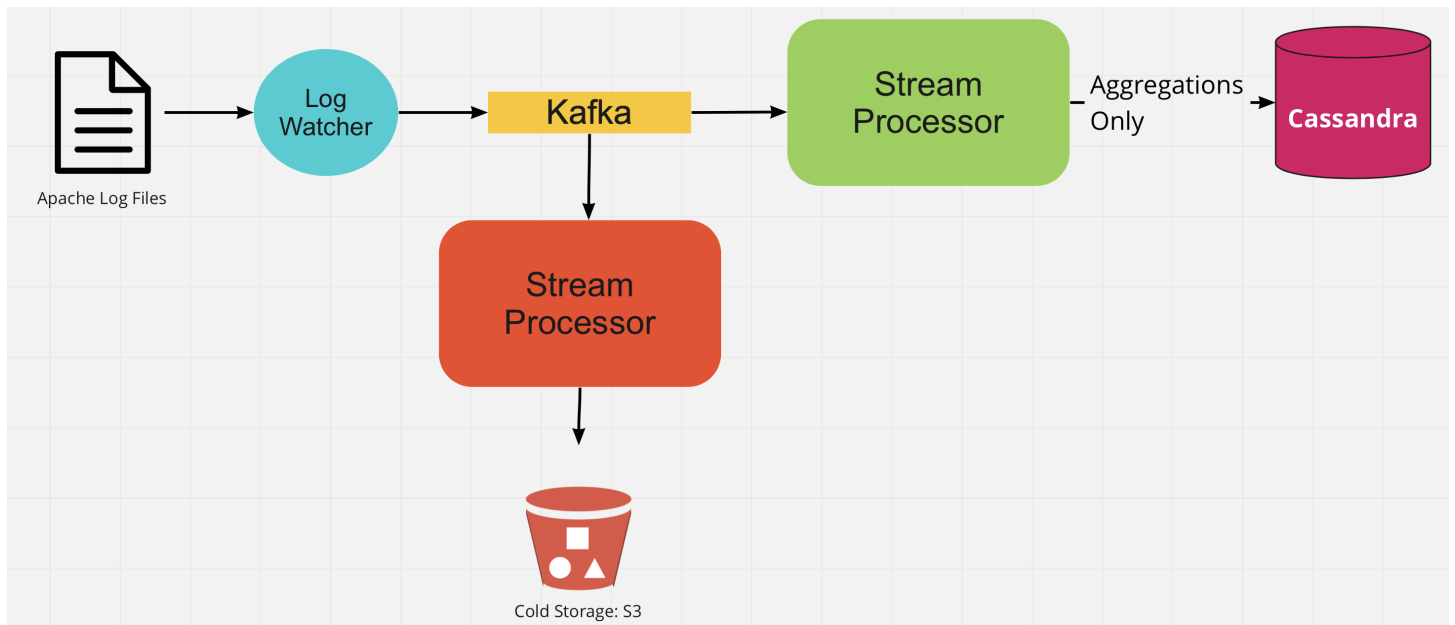
Issue 2: Data Model

Where do we store individual clicks?

Where do we store aggregations?

- Aggregation data:
 - Accessed frequently
 - Different time slices
- Individual clicks:
 - Rarely accessed
 - Historical / backfill / backup purposes

Let's not store it in our primary Cassandra table then.



Why use S3 for individual clicks?

- Cheaper to store
- Don't need to access often

How to store aggregations?

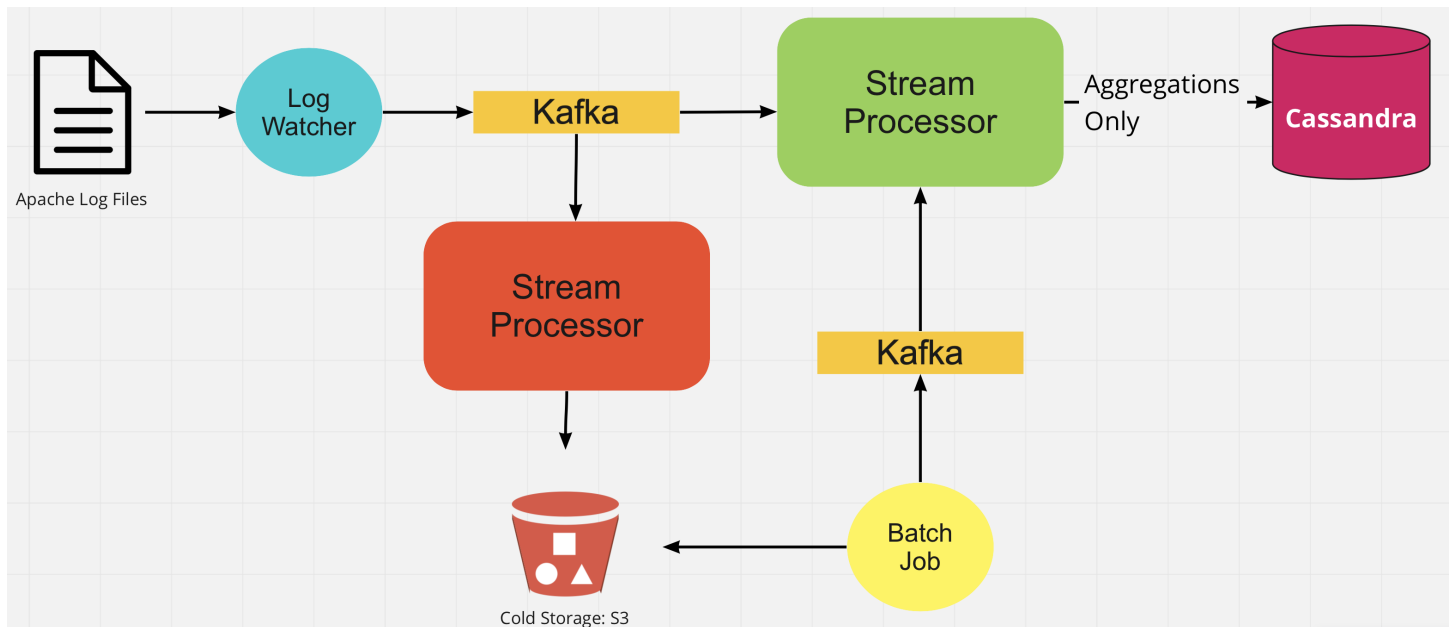
- Different time slices
- Redundant data is okay
- Design Cassandra tables based on query pattern

Issue 3: Data Inconsistency

Stream processors can miss **late data**

Aggregations can be slightly inaccurate

How can we fix that?



Scheduled batch job:

- Reads individual clicks from S3
- Re-publishes them to another Kafka topic
- Goes through stream processor (without late data now)
- Alternatively
 - Batch job directly calculates aggregations and updates Cassandra